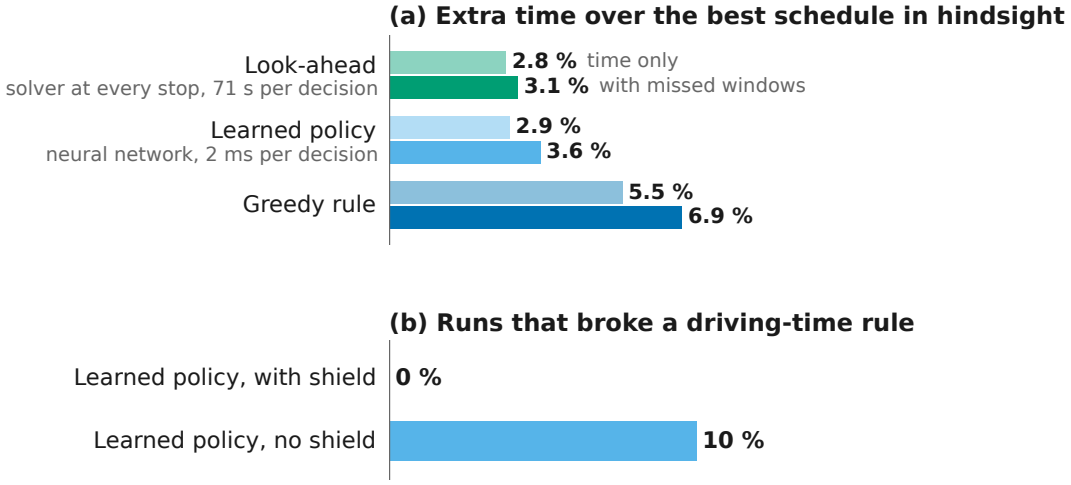
(a) Extra time over the best schedule in hindsight
Look-ahead
solver at every stop, 71 s per decision
2.8 % time only
3.1 % with missed windows
Learned policy
neural network, 2 ms per decision
2.9 %
3.6 %
Greedy rule
5.5 %
6.9 %
(b) Runs that broke a driving-time rule
Learned policy, with shield
0 %
Learned policy, no shield
10 %